

OpenAI API Complete Study Guide	2
1. How OpenAI Models Work.....	2
Core Architecture	2
Training Process	3
How Inference Works	3
Model Capabilities	3
Limitations	3
2. Setting Up a Request.....	3
Basic Setup	3
Essential Request Structure.....	4
Message Roles Explained	4
Common Configuration Options	4
Error Handling	5
3. Tokens.....	5
What Are Tokens?	5
Token Calculation	5
Token Economics	5
Managing Tokens	6
Token Strategies	6
4. Tools (Function Calling)	6
What Are Tools?.....	6
Setting Up Tools.....	6
Using Tools in Requests.....	7
Tool Use Cases	7
5. The "Few Shot" Approach	8
What Is Few-Shot Learning?.....	8
Basic Structure	8
Few-Shot Strategies	9
Best Practices	9
Pros and Cons	9
6. Temperature	9
What Is Temperature?	9

Temperature Values	10
Low Temperature (0 - 0.3)	10
Medium Temperature (0.4 - 1.2)	10
High Temperature (1.3 - 2.0).....	10
Temperature vs Top-p	11
7. Frequency and Presence Penalties	11
Presence Penalty	11
Frequency Penalty	11
Combined Usage	12
Practical Examples	12
Best Practices	12

OpenAI API Complete Study Guide

1. How OpenAI Models Work

Core Architecture

- **Transformer-based neural networks** trained on massive text datasets (trillions of tokens)

- **Autoregressive generation:** predicts the next token based on all previous tokens in sequence
- **Context window:** limited memory capacity (varies by model - GPT-3.5: ~4K tokens, GPT-4: up to 128K tokens)

Training Process

- **Pre-training:** learns patterns from internet text, books, articles
- **Fine-tuning:** specialized training on curated datasets
- **RLHF (Reinforcement Learning from Human Feedback):** aligned with human preferences

How Inference Works

- Input text is **tokenized** into numerical representations
- Model calculates **probability distributions** for next possible tokens
- **Sampling strategies** (temperature, top-p) determine which token is selected
- Process repeats until stopping condition is met

Model Capabilities

- **Text generation:** creative writing, summaries, explanations
- **Code generation:** programming in multiple languages
- **Analysis:** sentiment analysis, data interpretation
- **Conversation:** maintaining context across multi-turn dialogues

Limitations

- **Knowledge cutoff:** training data has a specific end date
- **Hallucinations:** can generate plausible but incorrect information
- **Context limits:** cannot remember beyond token window
- **Stateless:** each API call is independent

2. Setting Up a Request

Basic Setup

```
import OpenAI from "openai";
```

```
const openai = new OpenAI({
  apiKey: "your-api-key",
  dangerouslyAllowBrowser: true // Only for client-side
development
});
```

Essential Request Structure

```
const response = await openai.chat.completions.create({
  model: "gpt-4",
  messages: [
    {
      role: "system",
      content: "You are a helpful assistant."
    },
    {
      role: "user",
      content: "Hello, how are you?"
    }
  ]
});
```

Message Roles Explained

- **system:** Sets behavior and context (like giving instructions to the AI)
- **user:** Represents the human user's input
- **assistant:** The AI's responses (used for conversation history)
- **tool:** Used for function calling responses

Common Configuration Options

```
{
  model: "gpt-4",
  messages: messages,
  temperature: 1.0,           // Creativity level
  max_tokens: 150,           // Response length limit
  top_p: 1.0,                 // Nucleus sampling
  frequency_penalty: 0,       // Reduce repetition
  presence_penalty: 0,        // Encourage topic diversity
  stop: ["\n"],               // Stop sequences
}
```

```
    stream: false           // Real-time streaming
}
```

Error Handling

```
try {
    const response = await openai.chat.completions.create({...});
    console.log(response.choices[0].message.content);
} catch (error) {
    console.error("API Error:", error.message);
    // Handle rate limits, invalid requests, etc.
}
```

3. Tokens

What Are Tokens?

- **Not words or characters** - chunks of text with varying lengths
- **Average:** approximately 4 characters per token in English
- **Examples:**
 - "hello" = 1 token
 - "ChatGPT" = 2 tokens
 - "artificial intelligence" = 2 tokens

Token Calculation

- Use OpenAI's tokenizer: <https://platform.openai.com/tokenizer>
- **Prompt tokens:** input message tokens
- **Completion tokens:** generated response tokens
- **Total tokens:** prompt + completion

Token Economics

- **Cost:** charged per token (input + output)
- **Context limits:** total conversation must fit within model's token limit
- **Performance:** fewer tokens = faster responses

Managing Tokens

```
// Limiting output length
max_tokens: 100 // Maximum tokens in response

// Checking token usage
console.log(response.usage);
// Output: {prompt_tokens: 44, completion_tokens: 56, total_tokens: 100}
```

Token Strategies

- **Summarization:** compress long conversations to stay within limits
- **Chunking:** break large documents into token-sized pieces
- **Optimization:** remove unnecessary words and formatting

4. Tools (Function Calling)

What Are Tools?

- **Function calling:** allows models to call external functions
- **Structured outputs:** get JSON responses for specific tasks
- **Chain operations:** combine multiple API calls intelligently

Setting Up Tools

```
const tools = [
  {
    type: "function",
    function: {
      name: "get_stock_price",
      description: "Get current stock price for a symbol",
      parameters: {
        type: "object",
        properties: {
          symbol: {
            type: "string",
            description: "Stock symbol (e.g., AAPL,
TSLA)"
```

```

        }
      },
      required: ["symbol"]
    }
  }
}
];

```

Using Tools in Requests

```

const response = await openai.chat.completions.create({
  model: "gpt-4",
  messages: messages,
  tools: tools,
  tool_choice: "auto" // or "none", or specific function
});

// Handle function calls
if (response.choices[0].message.tool_calls) {
  const toolCall = response.choices[0].message.tool_calls[0];
  const functionName = toolCall.function.name;
  const functionArgs = JSON.parse(toolCall.function.arguments);

  // Execute your function
  const result = await executeFunction(functionName,
functionArgs);

  // Send result back to model
  messages.push({
    role: "tool",
    content: JSON.stringify(result),
    tool_call_id: toolCall.id
  });
}

```

Tool Use Cases

- **API integrations:** weather, stocks, databases
- **Calculations:** mathematical operations
- **Data retrieval:** file systems, web scraping

- **External services:** email, SMS, payments

5. The "Few Shot" Approach

What Is Few-Shot Learning?

- **Learning from examples:** providing input-output pairs as guidance
- **Pattern recognition:** model learns desired format and style
- **No training required:** works immediately with examples

Basic Structure

```
const messages = [  
  {  
    role: "system",  
    content: `You are a stock analyst. Format responses like the  
examples below:  
  
    ###  
    Stock: AAPL  
    Price: $150.25 (+2.3%)  
    Recommendation: BUY  
    Reason: Strong quarterly earnings  
    ###  
  
    Stock: MSFT  
    Price: $280.50 (-1.2%)  
    Recommendation: HOLD  
    Reason: Market volatility concerns  
    ###`  
  },  
  {  
    role: "user",  
    content: "Analyze TSLA stock"  
  }  
];
```


Few-Shot Strategies

- **Zero-shot:** no examples, just instructions
- **One-shot:** single example provided
- **Few-shot:** multiple examples (2-5 typically optimal)
- **Many-shot:** extensive examples (for complex tasks)

Best Practices

- **Clear delimiters:** use ### or --- to separate examples
- **Consistent format:** maintain same structure across examples
- **Diverse examples:** show different scenarios and edge cases
- **Quality over quantity:** better examples > more examples

Pros and Cons

Pros:

- More control over output format
- Consistent styling
- Works without fine-tuning
- Quick implementation

Cons:

- Uses more tokens (higher cost)
- Reduces available context space
- Can overfit to examples
- Less flexible than fine-tuning

6. Temperature

What Is Temperature?

- **Randomness control:** determines how "creative" or "conservative" responses are
- **Range:** 0 to 2 (default is 1)
- **Sampling parameter:** affects probability distribution of next token selection

Temperature Values

```
// Conservative (factual, predictable)
temperature: 0.2
```

```
// Balanced (default behavior)
temperature: 1.0
```

```
// Creative (diverse, unpredictable)
temperature: 1.8
```

Low Temperature (0 - 0.3)

- **Characteristics:** deterministic, factual, consistent
- **Use cases:**
 - Code generation
 - Mathematical calculations
 - Factual Q&A
 - Data analysis
 - Legal/medical content

Medium Temperature (0.4 - 1.2)

- **Characteristics:** balanced creativity and consistency
- **Use cases:**
 - General conversation
 - Content writing
 - Summaries
 - Explanations

High Temperature (1.3 - 2.0)

- **Characteristics:** highly creative, diverse, unpredictable
- **Use cases:**
 - Creative writing
 - Brainstorming
 - Poetry generation
 - Experimental content

Temperature vs Top-p

- **Alternative sampling methods:** use one or the other, not both at extremes
- **Temperature:** affects entire probability distribution
- **Top-p:** considers only top percentage of probable tokens

7. Frequency and Presence Penalties

Presence Penalty

What it does: Encourages discussing new topics **Range:** -2.0 to 2.0 (default: 0) **How it works:** Penalizes tokens that have already appeared

```
presence_penalty: 0.6 // Encourages topic diversity
```

Effects:

- **Low (0 to 0.5):** Focuses on current topics
- **Medium (0.5 to 1.0):** Balanced topic exploration
- **High (1.0 to 2.0):** Jumps between many topics

Use cases:

- **Low:** Detailed analysis of single topic
- **High:** Brainstorming multiple ideas

Frequency Penalty

What it does: Reduces repetition of the same phrases **Range:** -2.0 to 2.0 (default: 0)

How it works: Penalizes tokens based on frequency of use

```
frequency_penalty: 0.8 // Reduces repetitive language
```

Effects:

- **Low (0 to 0.5):** May repeat key phrases
- **Medium (0.5 to 1.0):** Varied language
- **High (1.0 to 2.0):** Avoids repetition strongly

Use cases:

- **Low:** Technical documentation (repetition needed)
- **High:** Creative writing (variety preferred)

Combined Usage

```
{
  temperature: 0.8,
  presence_penalty: 0.4,    // Slight topic diversity
  frequency_penalty: 0.6    // Reduce repetition
}
```

Practical Examples

For stock analysis:

```
// Focused, consistent analysis
{
  temperature: 0.3,
  presence_penalty: 0.1,
  frequency_penalty: 0.2
}
```

```
// Creative, varied reporting
{
  temperature: 1.1,
  presence_penalty: 0.6,
  frequency_penalty: 0.8
}
```

Best Practices

- **Start with defaults** (all at 0) and adjust incrementally
- **Test combinations** - penalties interact with temperature
- **Monitor output quality** - high penalties can hurt coherence
- **Use moderately** - extreme values (>1.5) often degrade performance
- **Consider your use case** - factual content needs different settings than creative content